

Step 1: Copy the Full Content

Copy everything from this block:

Kubernetes Assignments 1-5 Detailed Summary

Assignment 1: Deploy a Kubernetes Cluster & NGINX Deployment

Tasks

1. Deploy a Kubernetes cluster with 3 nodes.
2. Create an NGINX deployment with 3 replicas.

Deployment YAML

```
```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
 name: nginx-deployment
spec:
 replicas: 3
 selector:
 matchLabels:
 app: nginx
 template:
 metadata:
 labels:
 app: nginx
 spec:
 containers:
 - name: nginx
 image: nginx:latest
 ports:
 - containerPort: 80

```

**Commands**

```
kubectl apply -f nginx-deployment.yaml
kubectl get pods
```

## Problems Faced

- Error: unknown field "metadata.spec"
    - Cause: spec nested incorrectly inside metadata.
    - Solution: Correct YAML indentation.
  - Lesson: YAML indentation is critical in Kubernetes.
- 

## Assignment 2: Create NodePort Service for NGINX

### NodePort Service YAML

```
apiVersion: v1
kind: Service
metadata:
 name: nginx-nodeport
spec:
 type: NodePort # Exposes service outside cluster
 selector:
 app: nginx
 ports:
 - port: 80
 targetPort: 80
 nodePort: 30080
```

### Commands

```
kubectl apply -f nginx-nodeport.yaml
kubectl get svc
```

### Problems & Resolutions

1. Connection refused with curl
    - Cause: NodePort not directly accessible in WSL Minikube.
    - Solution: Use minikube service nginx-nodeport or port-forward.
  2. Minikube service "SVC\_NOT\_FOUND"
    - Cause: kubectl context was mycluster, not Minikube.
    - Solution: kubectl config use-context minikube
  3. Port already in use
    - Cause: Local port busy.
    - Solution: Use a different port, e.g., 8081.
-

## Assignment 3: Scale Deployment

### Command

```
kubectl scale deployment nginx-deployment --replicas=5
kubectl get pods
```

- Deployment automatically balances traffic across 5 pods.
  - Lesson: Services route traffic to all pods automatically.
- 

## Assignment 4: Change NodePort → ClusterIP

### ClusterIP Service YAML

```
apiVersion: v1
kind: Service
metadata:
 name: nginx-clusterip
spec:
```

type: ClusterIP # Internal-only access

selector:

app: nginx

ports:

- port: 80

targetPort: 80

- NodePort removed; service is internal.
  - Lesson: ClusterIP is default; NodePort exposes externally.
- 

## Assignment 5: Second NGINX Deployment + Ingress (Conceptual)


- Deploy another NGINX deployment (3 replicas).
  - Create ClusterIP service for new deployment.
  - Use Ingress to route paths: /app1 -> first NGINX, /app2 -> second NGINX.
  - Ingress acts as a single entry point for multiple services.
- 

## Key Learnings

1. YAML indentation matters.
2. Kubernetes contexts: Minikube vs other clusters.
3. NodePort may not work directly in WSL; use port-forward or minikube service.
4. Port-forward requires free local port.
5. Service types: ClusterIP = internal, NodePort = external, LoadBalancer = cloud LB.

6. Pods & service labels must match.
  7. Scaling replicas is simple; services automatically load-balance.
  8. Ingress centralizes traffic for multiple services.
- 

### Traffic Flow Diagram (Text Version)



### Summary Table of Service Types

Type	Accessible From	NodePort Needed?	Use Case
ClusterIP	Internal pods only	No	Internal communication between pods
NodePort	External + internal	Yes	Expose service outside cluster
LoadBalancer	External + internal	Optional	Cloud environments with LB

---

### Instructor Notes

- Always verify cluster, context, and namespace before creating resources.
- Use kubectl get pods, svc, endpoints for debugging.
- Service type determines accessibility.
- Replicas scaling is automatic via Deployment.
- Ingress routes multiple services via paths/domains.